

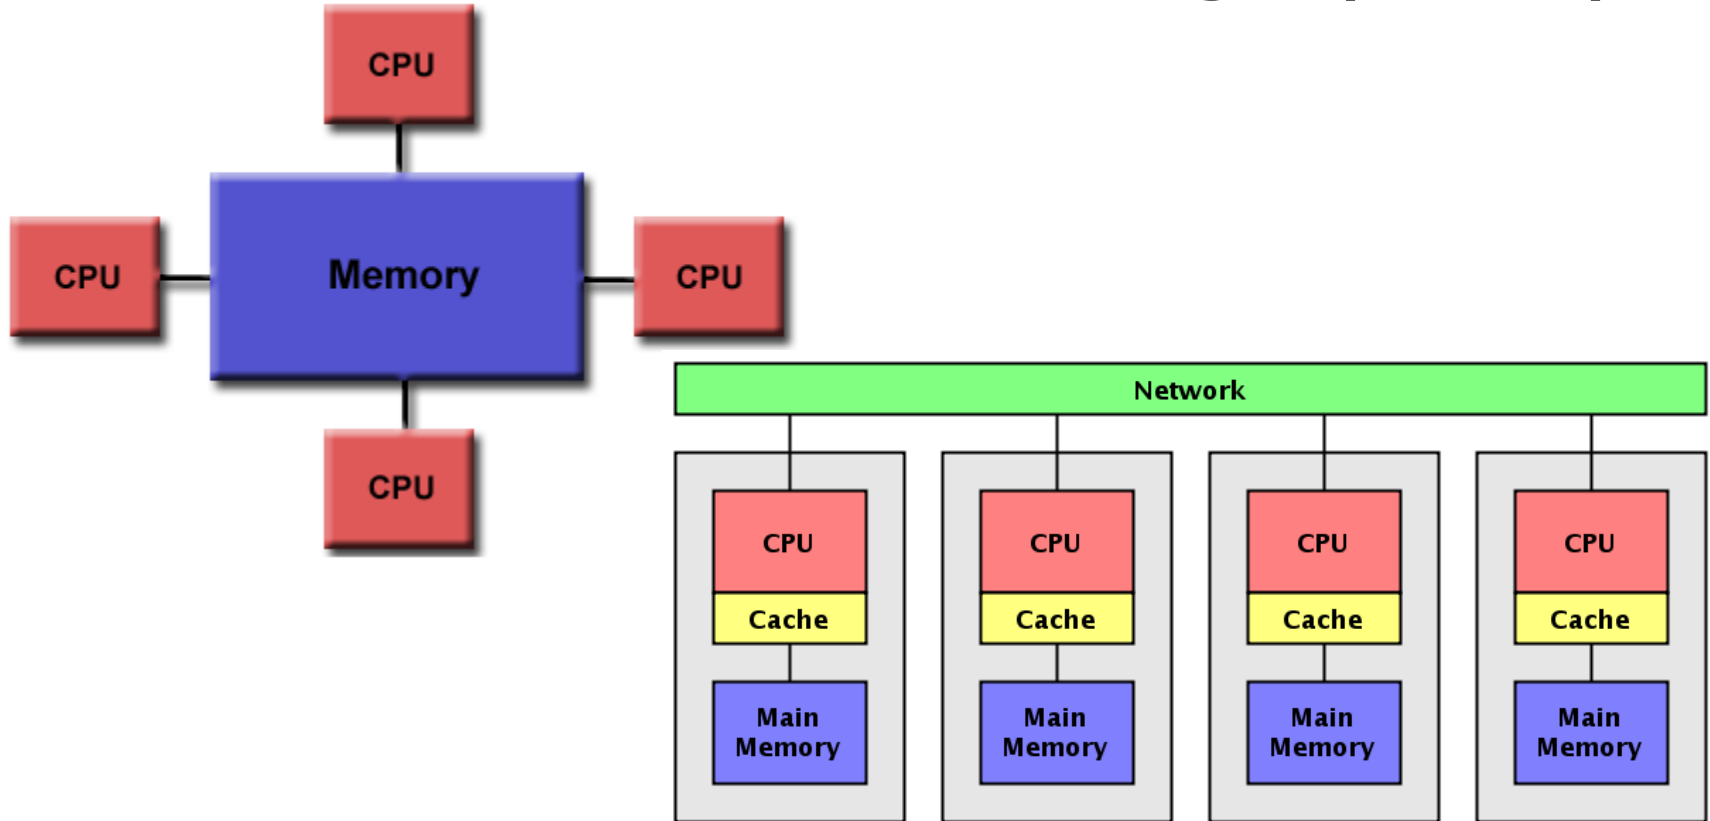
# Introduction to Parallel Computing

Michael Skuhersky  
vex@mit.edu

# What is Parallel Computing?

- Wikipedia says: “Parallel computing is a form of computation in which many calculations are carried out simultaneously”
- Speed measured in FLOPS

# What is Parallel Computing? (cont.)



# How can this be useful?

- Faster Calculations
- Make use of less powerful hardware?
- Many computers, or “nodes” can be combined into a cluster
- But, it's a lot more complex to implement

# Definitions

- Core: A processor that carries out instructions sequentially. Higher frequency means faster calculations
- Process/Task: A chain of instructions; a program

# Main Granularity Paradigms

- Granularity: the ratio of computation to communication
- 3 approaches to parallelism, depending on what kind of problem you need to solve



# Embarrassingly Parallel

- No effort required to separate tasks
- Tasks do not depend on, or communicate with, each other.
- Examples: Mandlebrot, Folding@home, Password brute-forcing, Bitcoin Mining!

# Example: Particle Physics



- Independent tracked particles
- They interact with the world, not each other
- Can be submitted as a batch of jobs





# Example: 3D Projection

- Each pixel on the screen, or each block of pixels, is rendered independently



# Coarse-grained Parallelism



- Tasks communicate with each other, but not more than once a second
- Examples: Stuff that involves synchronization

# Fine-grained Parallelism

- AKA Multithreading
- Subtasks must constantly communicate with each other
- Must use something like MPI

# Example: Molecular Dynamics

- Relaxation of a protein in water
- Movement of atoms depends on that of surrounding atoms



# Job Scheduling

- Integral to parallel computing; assigns tasks to cores
- Batch jobs, Multiple users, Resource sharing, System monitoring

# Livelock/Deadlock/Race Conditions

- Things that could go wrong when you are performing a fine or coarse-grained computation:
  - Livelock
  - Deadlock
  - Race Conditions


# Not Everything Benefits

- Many problems must be solved sequentially (e.g. Long division, protein folding)
- Interactive stuff



# Example: Protein Folding

- Each atom depends on the one before it



# Slowdown, Bottlenecks

- How fast can cores communicate with each other?
- How fast can nodes communicate with each other?
- Infiniband

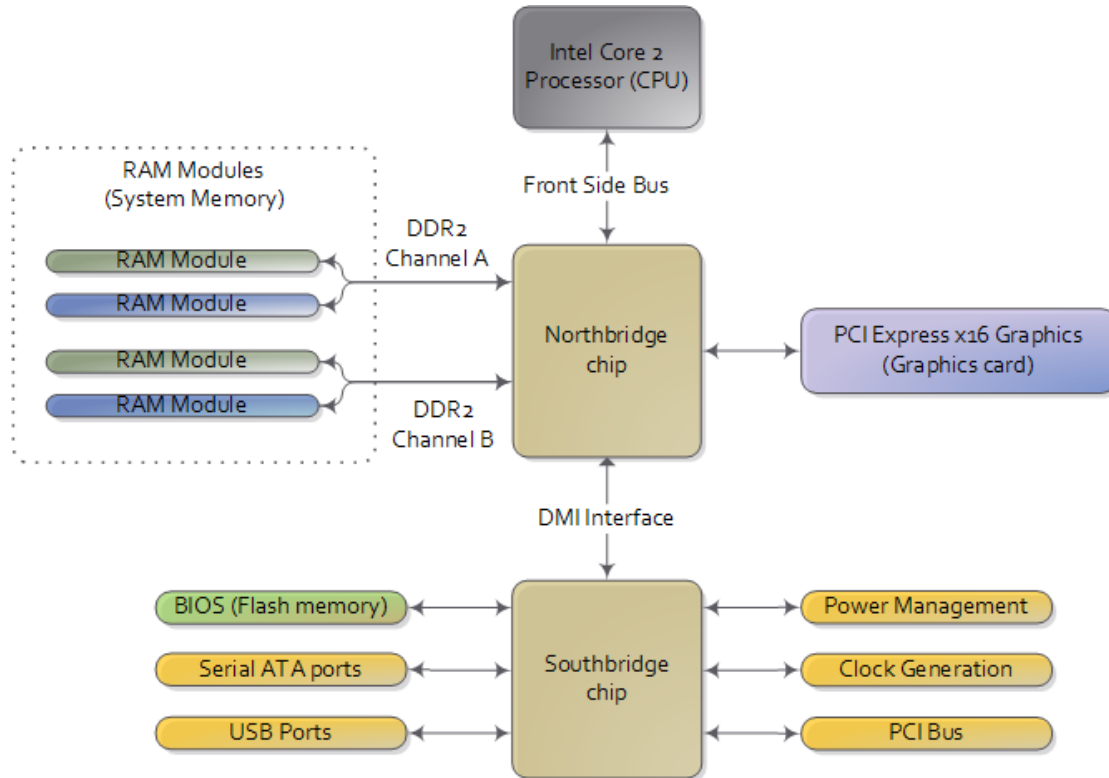
# Massively parallel computing

- Clusters, Supercomputers
- Grids (Folding@home, SETI@home, PlanetQuest)




# CPU Methods

# Computer Overview







CD-IN

SPDIF

1394 header

1394 controllers

Onboard LED

SATA RAID

Jumpers

SATA controller

USB header

Serial port connector

ATA controller

Game / MIDI header

FW in PLCC

System panel connectors

P4 power connector

Inductor (coil)

Capacitors

Gigabit LAN header

Osculator

CPU socket lever

CPU socket

Voltage regulator

Northbridge

4-pin CPU fan connection

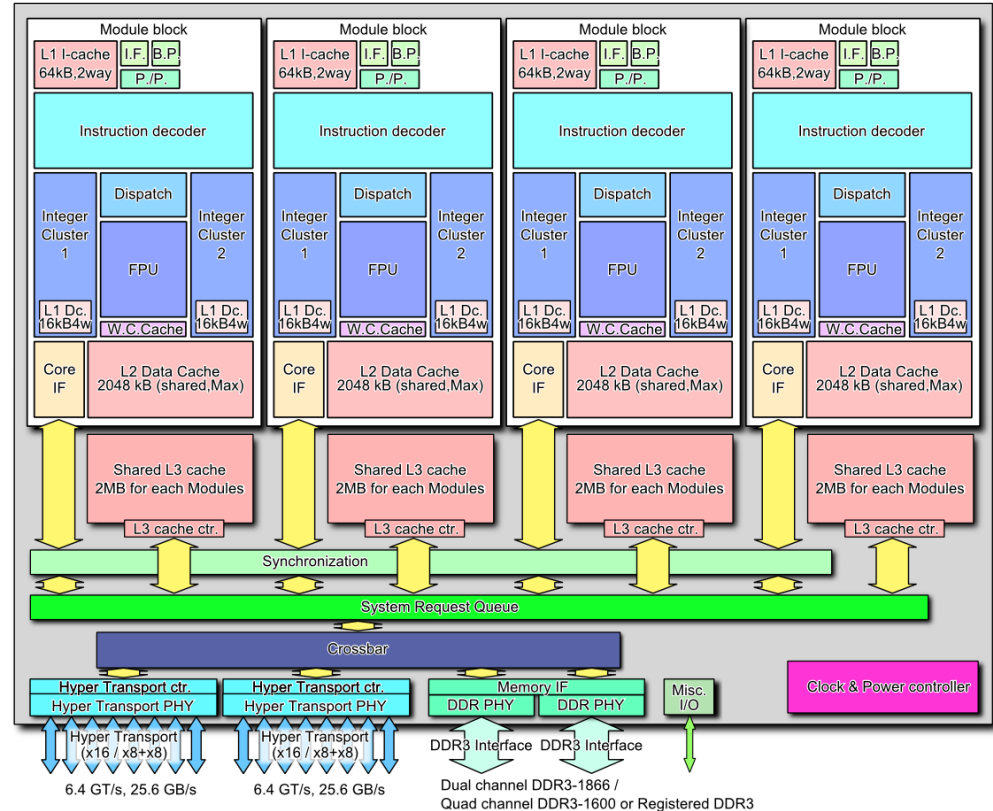


Motherboard model name

4x DIMM memory slots

Super I/O



# CPU Architecture





# **CPUs can't get much faster!**

- Temperature
- Lithography limitations
- Quantum tunneling
- Electricity travel speed
- We can add more cores though

# Message Passing Interface

- Allows individual processes to talk to processes on different cores

# MPI Implementations

- OpenMPI, MPICH
- Wrappers for various languages
- `mpirun -np process-count program-name`

# Example Code


```
#define NPTS 1000000000
#define NLOOP 10

void main(int argc, char** argv) {
    int rank,nproc,i,istart,iend,loop,N;
    unsigned long int start_time,end_time;
    struct timeval start,end;
    double sum,pi,mflops;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &nproc);
    if (rank == 0) {
        gettimeofday(&start, NULL);
        istart = 1 + NPTS*((rank+0.0)/nproc);
        iend = NPTS*((rank+1.0)/nproc);
        sum = 0.0;
        for (loop = 0; loop < NLOOP; ++loop)
            for (i = istart; i <= iend; ++i)
                sum += 0.5/((i-0.75)*(i-0.25));
        MPI_Reduce(&sum,&pi,1,MPI_DOUBLE,MPI_SUM,0,MPI_COMM_WORLD);
        gettimeofday(&end, NULL);
        start_time = start.tv_sec * 1e6 + start.tv_usec;
        end_time = end.tv_sec * 1e6 + end.tv_usec;
        mflops = NLOOP*(NPTS*(5.0/(end_time-start_time)));
        printf("processes = %d, NPTS = %d, NLOOP = %d, pi = %f\n",nproc,NPTS,NLOOP,
pi/NLOOP);
        printf("time = %f, estimated MFlops = %f\n", (end_time-start_time)/1.0e6,mflops);
    }
    else {
        istart = 1 + NPTS*((rank+0.0)/nproc);
        iend = NPTS*((rank+1.0)/nproc);
        sum = 0.0;
        for (loop = 0; loop < NLOOP; ++loop)
            for (i = istart; i <= iend; ++i)
                sum += 0.5/((i-0.75)*(i-0.25));
        MPI_Reduce(&sum,&pi,1,MPI_DOUBLE,MPI_SUM,0,MPI_COMM_WORLD);
    }
    MPI_Finalize();
}
```

# Expanding Into a Cluster

- Multiple computers can be linked together over a network. MPI can use this



# GPU Methods

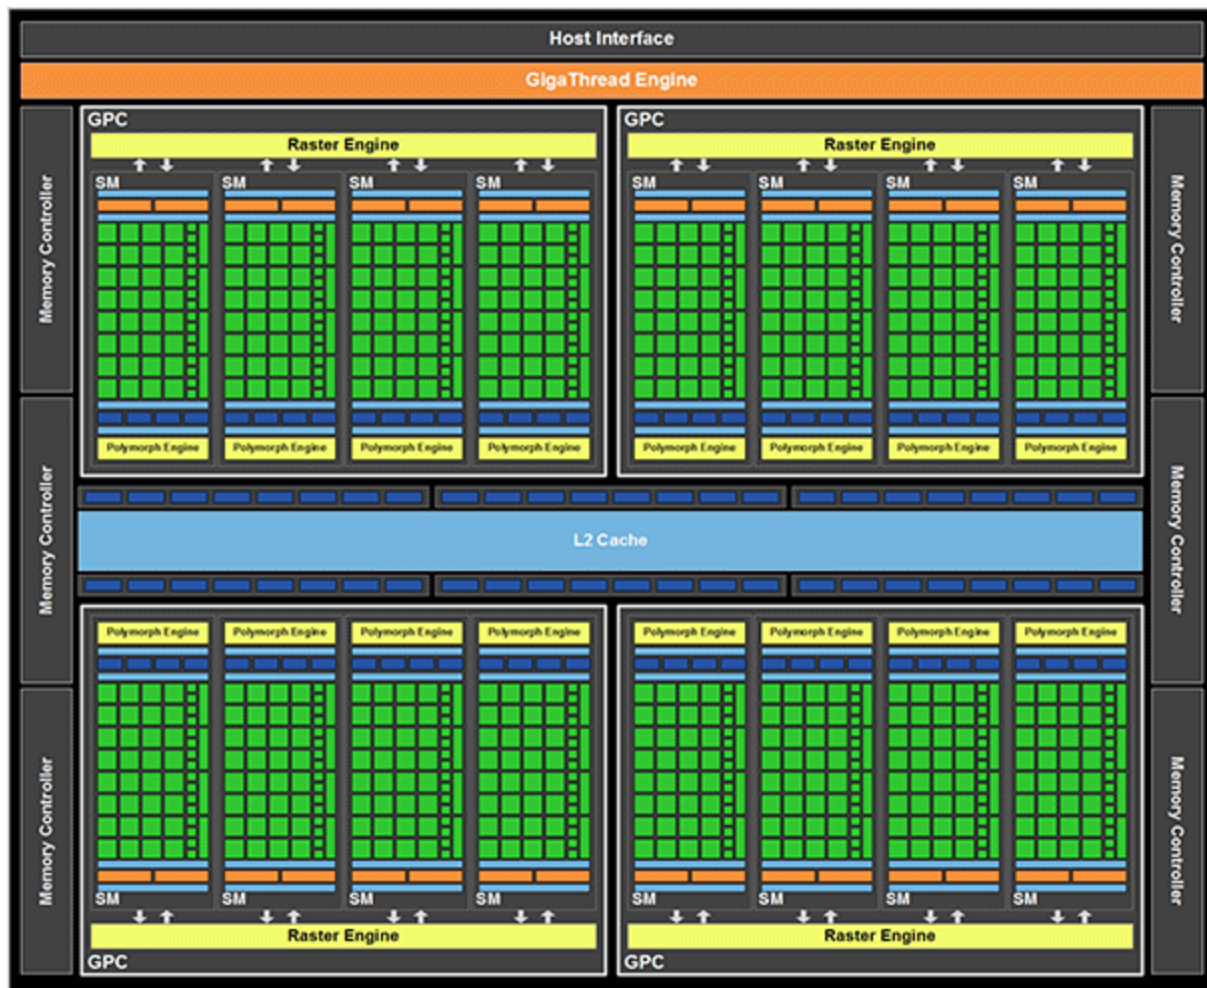
# Why were GPUs created?

- Simply, wanted to free up CPU
- GUIs required programmers to think in different ways
- In a GUI, everything behaves independently

# GPU Architecture

- Like a multi-core CPU, but with thousands of cores
- Has its own memory to calculate with





# GPU Advantages


- Ridiculously higher net computation power than CPUs
- Can be thousands of simultaneous calculations
- Pretty cheap

# Historic GPU Programming

- First developed to copy bitmaps around
- OpenGL, DirectX
- These APIs simplified making 3D games/visualizations

# Pipeline for rendering 3D

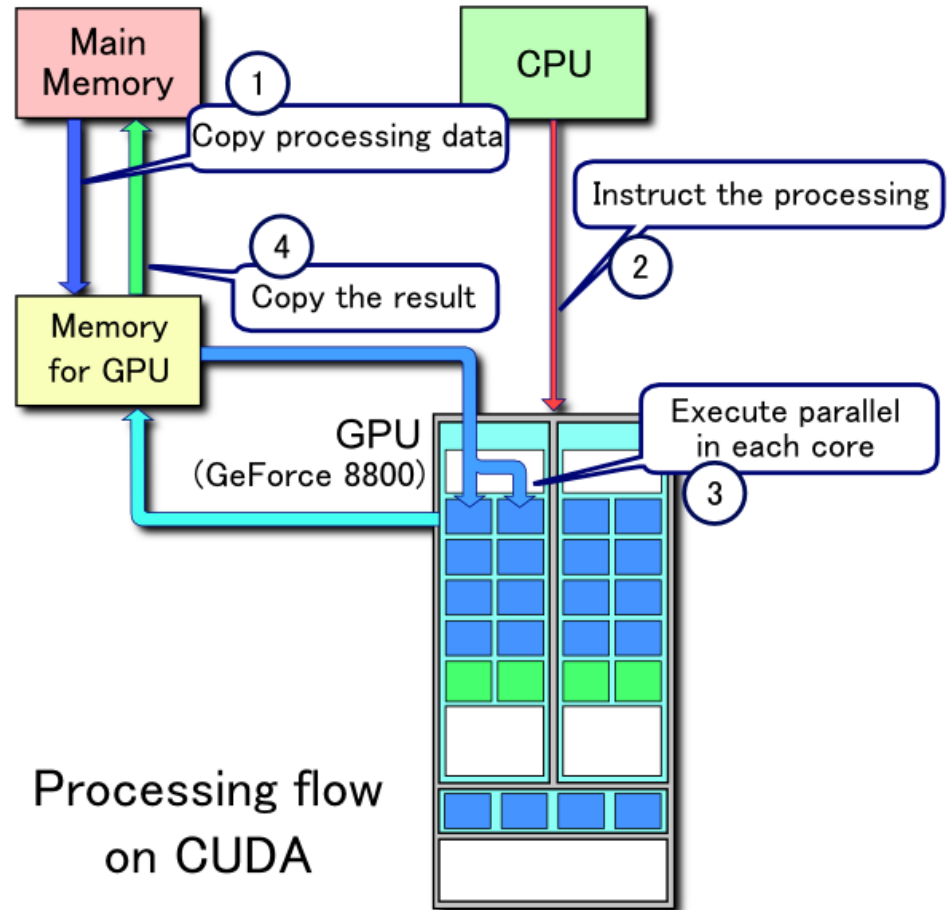
- Vertex data sent in by graphics API  
(from CPU code via OpenGL or DirectX, for example)
- Processed by vertex program (shader)
- Rasterized into pixels
- Processed by fragment shader



# Modern GPU Frameworks

- CUDA: Proprietary, Easy to use, Sponsored by NVIDIA and only runs on their cards
- OpenCL: Open, a bit harder to use, runs on both ATI and NVIDIA arch

# CUDA Methods



# CUDA Example Code

## Standard C Code

```
void saxpy(int n, float a,
          float *x, float *y)
{
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}

int N = 1<<20;

// Perform SAXPY on 1M elements
saxpy(N, 2.0, x, y);
```

## C with CUDA extensions

```
__global__
void saxpy(int n, float a,
          float *x, float *y)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < n) y[i] = a*x[i] + y[i];
}

int N = 1<<20;
cudaMemcpy(x, d_x, N, cudaMemcpyHostToDevice);
cudaMemcpy(y, d_y, N, cudaMemcpyHostToDevice);

// Perform SAXPY on 1M elements
saxpy<<<4096,256>>>(N, 2.0, x, y);

cudaMemcpy(d_y, y, N, cudaMemcpyDeviceToHost);
```

```
//Vector size in elements
const int N = 1048576;
//Vector size in bytes
const int dataSize = N * sizeof(float);
```

```
//CPU memory allocation
float *h_A = (float *)malloc(dataSize);
float *h_B = (float *)malloc(dataSize);
float *h_C = (float *)malloc(dataSize);
```

```
//GPU memory allocation
float *d_A, *d_B, *d_C;
cudaMalloc((void **)&d_A, dataSize);
cudaMalloc((void **)&d_B, dataSize);
cudaMalloc((void **)&d_C, dataSize);
```

```
//Initialize h_A[], h_B[]...
```

```
//Copy input data to GPU for processing
cudaMemcpy(d_A, h_A, dataSize, cudaMemcpyHostToDevice);
cudaMemcpy(d_B, h_B, dataSize, cudaMemcpyHostToDevice);
```

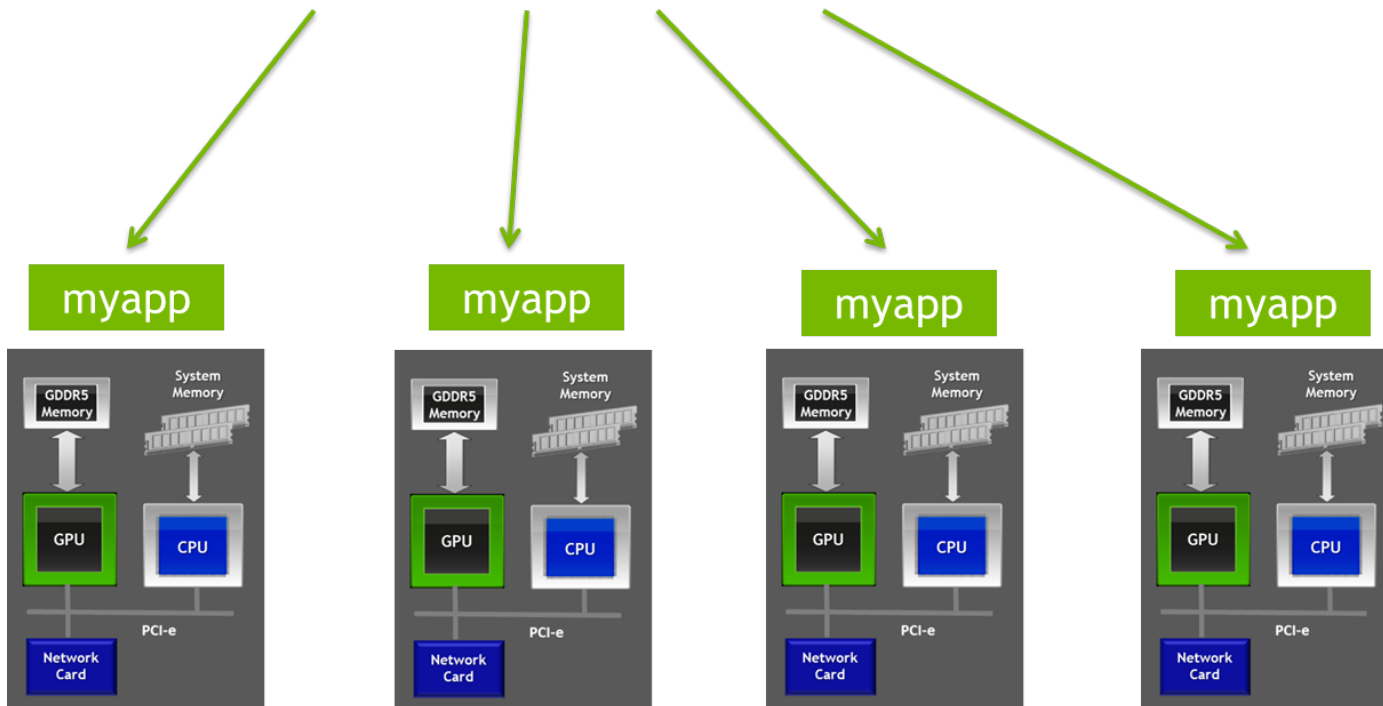
```
//Run the core of N / 256 units, 256 streams each
//Assuming that N is multiple of 256
vectorAdd<<<N / 256, 256>>>(d_C, d_A, d_B);
```

```
//Read GPU results
cudaMemcpy(h_C, d_C, dataSize, cudaMemcpyDeviceToHost);
```



# Can be used with MPI

```
mpirun -np 4 ./myapp <args>
```



# **How to Setup Your Own Parallel Computer**

# CPU or GPU Based?

- CPU: Easier to program for, has much more powerful individual cores
- GPU: Trickier to program for, thousands of really weak cores

# Cluster or Multicore?

- Multicore: All cores on in a single computer, usually shared memory.
- Cluster: Many computers linked together, each with individual memory

# OS/Distribution?

- Updated or Stable?
- Linux or something crazy?
- Custom cluster distro?

# Job Scheduler?



- Submit stuff in batches or MPI?
- Grid Engine, Condor, SLURM

# Programming It

# **Other Methods and their Applications**



# FPGAs



# Cryptography

- numbers
- brute-forcing

# Bitcoin

- Proof-Of Work